

PRODUCT DEMAND ANALYSIS

PROBLEM STATEMENT:

Create a machine learning model that forecasts product demand based on historical sales and external factors, helping businesses optimize inventory management and production planning to meet customer needs efficiently.

PHASES OF DEVELOPMENT:

PHASE 1: PROBLEM DEFINITION AND DESIGN THINKING

PROBLEM DEFINITION:

The problem is to create a machine learning model that forecasts product demand based on historical sales data and external factors. The goal is to help businesses optimize inventory management and production planning to efficiently meet customer needs. This project involves data collection, data preprocessing, feature engineering, model selection, training and evaluation.

DESIGN THINKING:

1. **Data Collection:** Collect historical sales data and external factors that influence demand, such as marketing campaigns, holidays, economic indicators, etc.
2. **Data Preprocessing:** Clean and preprocess the data, handle missing values, and convert categorical features into numerical representations.
3. **Feature Engineering:** Create additional features that capture seasonal patterns, trends, and external influences on product demand.
4. **Model Selection:** Choose suitable regression algorithms (e.g., Linear Regression, Random Forest, XG Boost) for demand forecasting.
5. **Model Training:** Train the selected model using the preprocessed data.
6. **Evaluation:** Evaluate the model's performance using appropriate regression metrics (e.g., Mean Absolute Error, Root Mean Squared Error).

PHASE 2: INNOVATION

By using incorporating time series forecasting techniques like ARIMA or Prophet for capture temporal patterns in demand data.

DATA SET LINK: <https://www.kaggle.com/datasets/chakradharmattapalli/product-demand-prediction-with-machine-learning>

PHASE 3: DEVELOPMENT PART 1

Begin building the product demand prediction model by loading and preprocessing the dataset.

PHASE 4: DEVELOPMENT PART 2

Continue building the product demand prediction model by feature engineering, model training, and evaluation.

PHASE 5: PROJECT DOCUMENTATION & SUBMISSION

Document the product demand prediction project and prepare it for submission.

DESCRIPTION OF THE DATASET:

The dataset contain the index of store id, product id, total price, base price and unsold units for 4 columns and contain the row 150151 products.

DATA PREPROCESSING STEPS:

1. Import all the necessary libraries
2. Import the dataset and read it as CSV file.
3. Check the null values
4. Fill the NA values

ANALYSIS TECHNIQUES APPLIED:

For analysis time series forecasting techniques like ARIMA model is used to capture temporal patterns in demand data. In the **ARIMA** model **PACF** and **TSEA** plots are used.

After analysis with ARIMA model perform the encoding categorical data, feature engineering (selection), splitting the data for model training, model training and evaluation.

In the model training there are several model like ‘**Linear Regression, Random forest, Support vector machine and Gradient Boosting**’ are used in predicting the model.

In model the ‘Mean Squared Error, R-squared’ are used to evaluating the model.

PROGRAM:

```
# Import necessary libraries
```

```
import pandas as pd
```

```
from sklearn.model_selection import train_test_split
```

```
from sklearn.preprocessing import StandardScaler
```

```
# Data Collection
```

```
# Assuming your dataset is named 'product_demand_data.csv' and located in the same directory  
as your Python script
```

```
data = pd.read_csv('/content/ProductDemand.csv')
```

```
# Data Preprocessing
```

```
# Handling Missing Values (if any)
```

```
data.fillna(0, inplace=True)
```

```
data.isnull().sum()
```

```
# Data Transformation
```

```
# No categorical variables to encode in this case
```

```
# Split Data
```

```
X = data[features] # Features
```

```
y = data[target] # Target variable
```

```
# Split the data into training and testing sets (70-30 split)
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)

# Data Standardization (optional, but often necessary for many machine learning algorithms)
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)

# Encode categorical data

import numpy as np
import pandas as pd

# One-hot encode the categorical data
encoded_df = pd.get_dummies(data)
# Print the encoded DataFrame
print(encoded_df)

# Feature Selection
features = ['ID', 'Store ID', 'Total Price', 'Base Price'] # Features
target = 'Units Sold' # Target variable

# Histograms and Box Plots

import matplotlib.pyplot as plt

# Histograms
data[features].hist(bins=20, figsize=(12, 10))
plt.suptitle("Histograms of Features")
```

```
plt.show()
```

```
# Box Plots
```

```
data[features].plot(kind='box', vert=False, figsize=(12, 6))
```

```
plt.title("Box Plots of Features")
```

```
plt.show()
```

```
# Correlation Matrix
```

```
import seaborn as sns
```

```
correlation_matrix = data[features].corr()
```

```
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm')
```

```
plt.title("Correlation Matrix")
```

```
plt.show()
```

```
# Pair Plot
```

```
sns.pairplot(data[features])
```

```
plt.suptitle("Pair Plot of Features")
```

```
plt.show()
```

```
# Target Variable Distribution
```

```
plt.figure(figsize=(8, 6))
```

```
sns.histplot(data[target], bins=20, kde=True)
```

```
plt.title("Distribution of Target Variable")
```

```
plt.xlabel(target)
```

```
plt.ylabel("Frequency")
```

```
plt.show()
```

```
# Feature vs. Target Plots
```

```
for feature in features:
```

```
    plt.figure(figsize=(8, 6))
```

```
    sns.scatterplot(x=data[feature], y=data[target])
```

```
    plt.title(f"{feature} vs. {target}")
```

```
    plt.xlabel(feature)
```

```
    plt.ylabel(target)
```

```
    plt.show()
```

```
# Box Plot of Target Variable Grouped by Categorical Feature
```

```
categorical_feature = 'Store ID' # Example categorical feature
```

```
plt.figure(figsize=(10, 6))
```

```
sns.boxplot(x=categorical_feature, y=target, data=data)
```

```
plt.title(f"Box Plot of {target} Grouped by {categorical_feature}")
```

```
plt.xlabel(categorical_feature)
```

```
plt.ylabel(target)
```

```
plt.xticks(rotation=45)
```

```
plt.show()
```

MODEL SELECTION

```
# Import necessary libraries for different algorithms
```

```
from sklearn.linear_model import LinearRegression
```

```
from sklearn.ensemble import RandomForestRegressor, GradientBoostingRegressor
```

```
from sklearn.svm import SVR

from sklearn.metrics import mean_squared_error, r2_score


# Initialize models

linear_reg = LinearRegression()

random_forest = RandomForestRegressor(random_state=42)

svm = SVR()

gradient_boosting = GradientBoostingRegressor(random_state=42)


# Train and predict using each algorithm

models = [linear_reg, random_forest, svm, gradient_boosting]

model_names = ['Linear Regression', 'Random Forest', 'Support Vector Machine', 'Gradient Boosting']

for model, name in zip(models, model_names):

    model.fit(X_train, y_train)

    predictions = model.predict(X_test)

    mse = mean_squared_error(y_test, predictions)

    r2 = r2_score(y_test, predictions)

    print(f"Model: {name}")

    print(f"Mean Squared Error: {mse:.2f}")

    print(f"R-squared: {r2:.2f}")

    print("-" * 30)
```

OUTPUT:

Data Preprocessing

```
ID          0
Store ID    0
Total Price 0
Base Price  0
Units Sold  0
dtype: int64
```





SPLIT DATA

Encoding categorical data


| | ID | Store ID | Total Price | Base Price | Units Sold |
|--------|--------|----------|-------------|------------|------------|
| 0 | 1 | 8091 | 99.0375 | 111.8625 | 20 |
| 1 | 2 | 8091 | 99.0375 | 99.0375 | 28 |
| 2 | 3 | 8091 | 133.9500 | 133.9500 | 19 |
| 3 | 4 | 8091 | 133.9500 | 133.9500 | 44 |
| 4 | 5 | 8091 | 141.0750 | 141.0750 | 52 |
| ... | ... | ... | ... | ... | ... |
| 150145 | 212638 | 9984 | 235.8375 | 235.8375 | 38 |
| 150146 | 212639 | 9984 | 235.8375 | 235.8375 | 30 |
| 150147 | 212642 | 9984 | 357.6750 | 483.7875 | 31 |
| 150148 | 212643 | 9984 | 141.7875 | 191.6625 | 12 |
| 150149 | 212644 | 9984 | 234.4125 | 234.4125 | 15 |

[150150 rows x 5 columns]

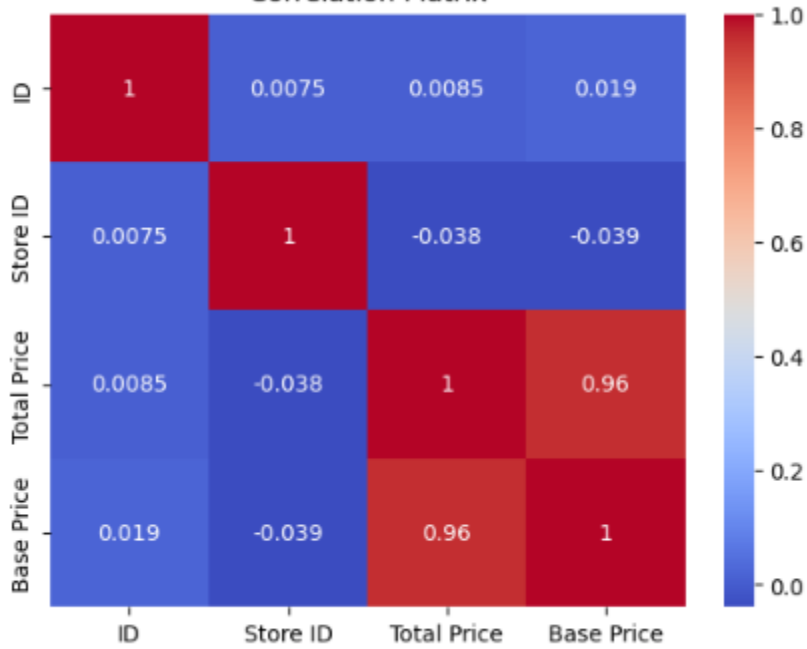
Histograms of Features

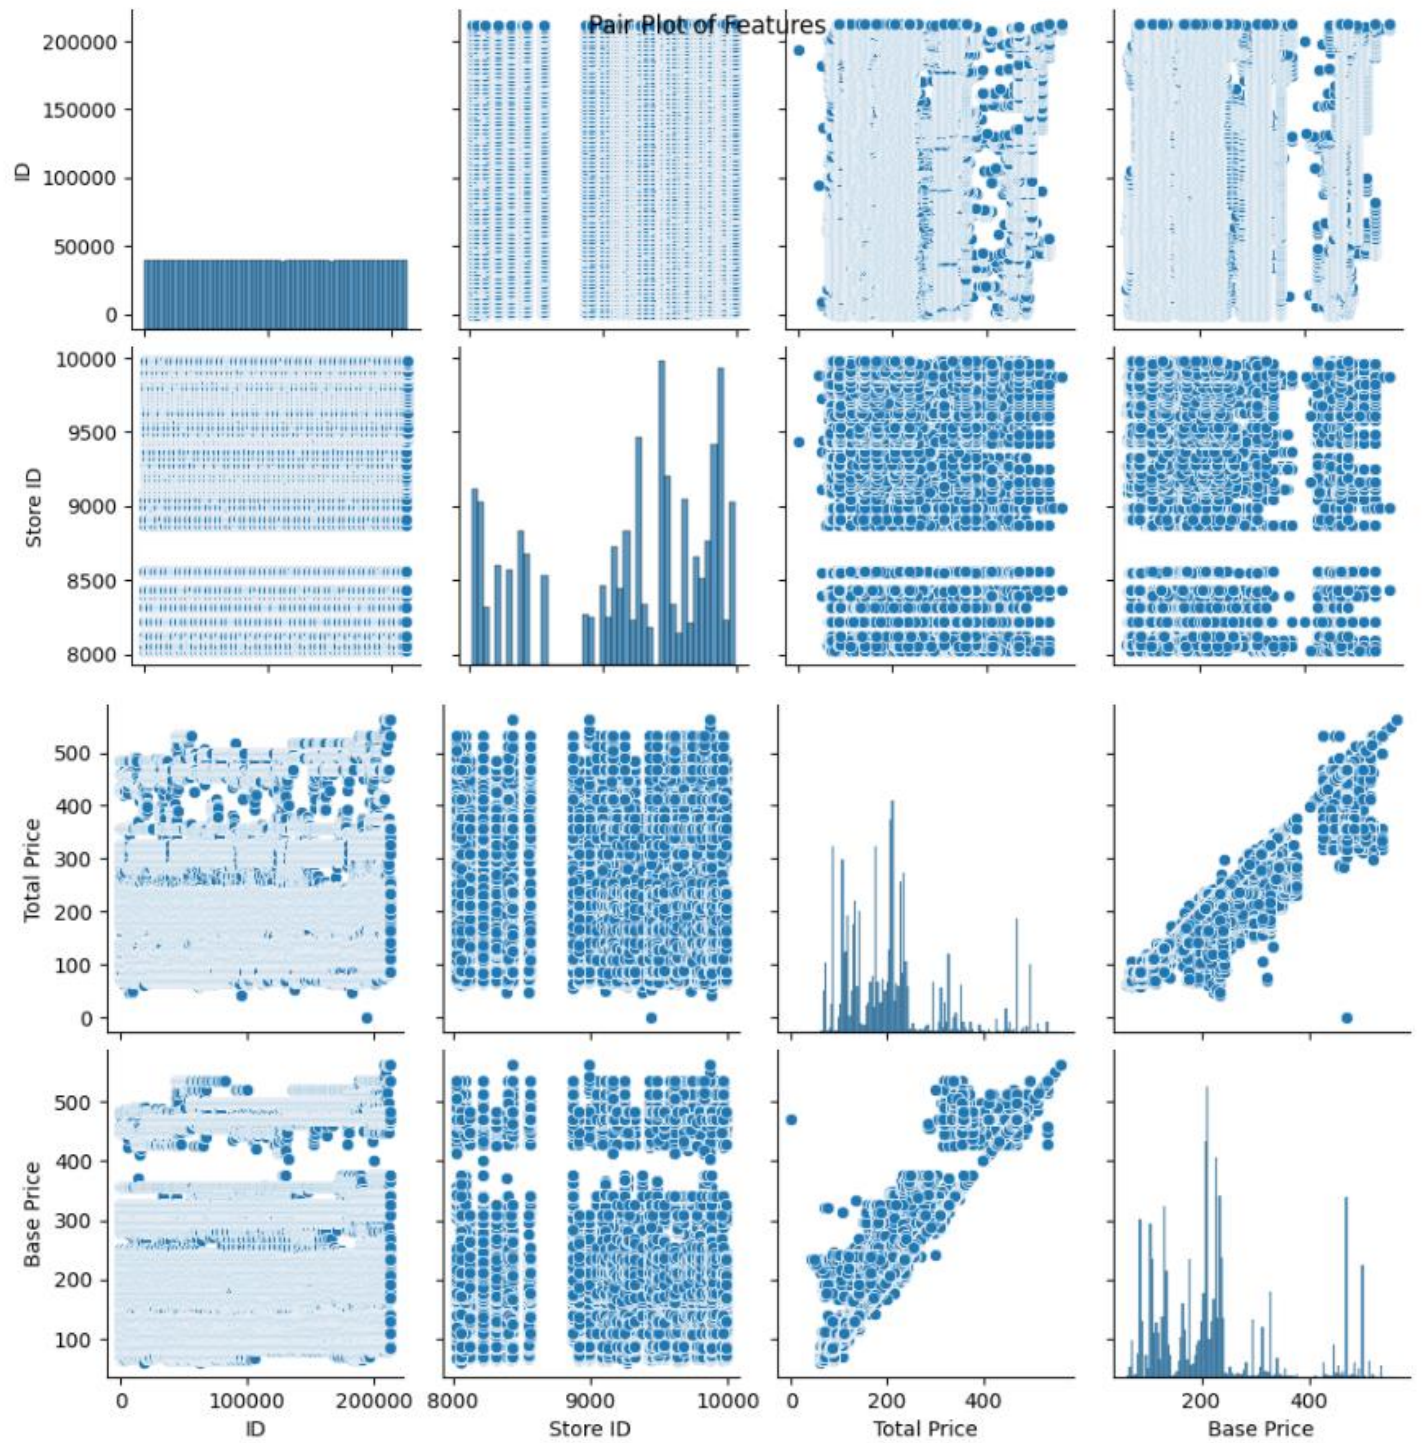


Box Plots of Features




Correlation Matrix







ID vs. Units Sold




Distribution of Target Variable




Store ID vs. Units Sold




Total Price vs. Units Sold



Base Price vs. Units Sold



Box Plot of Units Sold Grouped by Store ID



```
Model: Linear Regression  
Mean Squared Error: 2844.00  
R-squared: 0.15
```

```
-----  
Model: Random Forest  
Mean Squared Error: 1156.38  
R-squared: 0.66
```

```
-----  
Model: Support Vector Machine  
Mean Squared Error: 2956.17  
R-squared: 0.12
```

```
-----  
Model: Gradient Boosting  
Mean Squared Error: 1885.63  
R-squared: 0.44  
-----
```

ARIMA MODEL:

LINK:

<https://colab.research.google.com/drive/1gNQWTceGFUHsXDgdAMRgesCE5hFBYAAat?usp=sharing>

CONCLUSION:

Random Forest and Gradient Boosting model typically perform well in a variety of datasets due to their ability to capture complex patterns in the data.

Support Vector Machine (SVM) might perform well if the dataset has high dimensionality and complex relationships, although it might require fine-tuning of hyper parameter for optimal results.

Linear Regression provides a basic understanding of the relationships between variables but might not capture intricate patterns present in the data.